

Adaptive Foveated 2.5D Semantic Mapping for Autonomous Navigation

High-Speed Variable-Resolution Perception & Real-Time Kinematic Control for Tactical UGVs

Problem Statement Details

- Problem Statement ID: 26053
- Category / Domain: Software | Ministry of Defence / DRDO
- Theme: Smart Vehicles / Autonomous Systems
- Target Deployment: Defense Unmanned Ground Vehicles (UGVs)
- Operational Scope: Off-road, rough terrain, GPS-denied environments
- Core Breakthrough: Foveated 2.5D variable-resolution spatial hashing

Team TeraPulse — IIIT Bhopal

- Team Leader: Abhay Verma (Pipeline Integration & Control)
- Member 2: Lokendra Singh (Adaptive Grid Engine & Benchmark)
- Member 3: Yashvardhan Jain (LiDAR Preprocessing & Noise Rejection)
- Member 4: Chetan Meena (Semantic Terrain Segmentation & ML)
- Member 5: Sachin Chaubey (Visualization, UI Dashboard & Testing)
- Member 6: Esha Verma (Dataset Curation, QA & Documentation)
- Institute: Indian Institute of Information Technology Bhopal
- Repository: github.com/AbhayVerma628/SIH26053_Lidar

❖ Proposed Solution: TERA PULSE Architecture

- Foveated 2.5D Elevation Grid: Transforms millions of 3D LiDAR points into an egocentric multi-resolution heightmap anchored to the vehicle.
- Near Field (0 – 10 m): Ultra-fine 0.25 m resolution for millimetric obstacle clearance, step height, and reactive steering.
- Middle Field (10 – 30 m): Medium 0.75 m resolution capturing corridor contours and upcoming turns without memory bloat.
- Far Field (30 – 100 m): Coarse 2.00 m resolution providing global situational context with minimal cell allocations.
- Semantic Terrain Profiling: Every cell stores Z_{max} , Z_{mean} , slope gradient, surface roughness, and classification (Ground, Obstacle, Dynamic).
- Closed-Loop Navigation: Feeds directly into A* path planning and Pure Pursuit kinematics for autonomous vehicle motion.

❖ Innovation & Uniqueness

- Biomimetic Human-Eye Vision: Replaces uniform voxel brute-force with concentric attention bands, slashing compute load by 66.2%.
- 2.5D vs 3D Advantage: Delivers 90% of 3D geometric intelligence (slope, elevation, roughness) at the speed and lightweight RAM of a 2D map.
- Dynamic Moving Resolution Bubble: As the car drives, the fine 0.25 m bubble continuously moves with the vehicle $[x(t), y(t)]$ at 60 FPS.
- Vertical Beam Clearance Analysis: Distinguishes solid walls from hollow underpasses/tunnels without requiring heavy 3D voxelization.
- Edge Deployment Ready: Runs in 24.23 ms (>35 FPS), allowing real-time deployment on rugged low-power NVIDIA Jetson Orin & Raspberry Pi.

66.2% Cell Reduction

2,888 uniform cells down to 976 adaptive cells

24.23 ms Mapping Latency

4.8× faster than uniform grid (115.9 ms)

83.7% Drivable Space

Accurately separated safe corridor from obstacles

End-to-End Autonomous Pipeline

► Stage 1: Preprocessing & Filtering

Voxel grid downsampling (0.15m), Statistical Outlier Removal (SOR, 20 neighbors), RANSAC ground plane extraction.

► Stage 2: Deep Semantic Classification

Terrain classified into 4 functional classes: Ground (0), Static Obstacle (1), Dynamic Vehicle/Object (2), and Unknown (3).

► Stage 3: Adaptive 2.5D Grid Engine

Euclidean spatial hashing allocates points into concentric 0.25m / 0.75m / 2.0m cells, calculating Z_{max} , Z_{mean} , slope, and roughness.

► Stage 4: Traversability Cost Modeling

Multi-factor traversability cost: $Cost = w_{slope} * (slope) + w_{rough} * (roughness) + w_{obs} * (obstacle_cost)$. Safe threshold filters drivable ground.

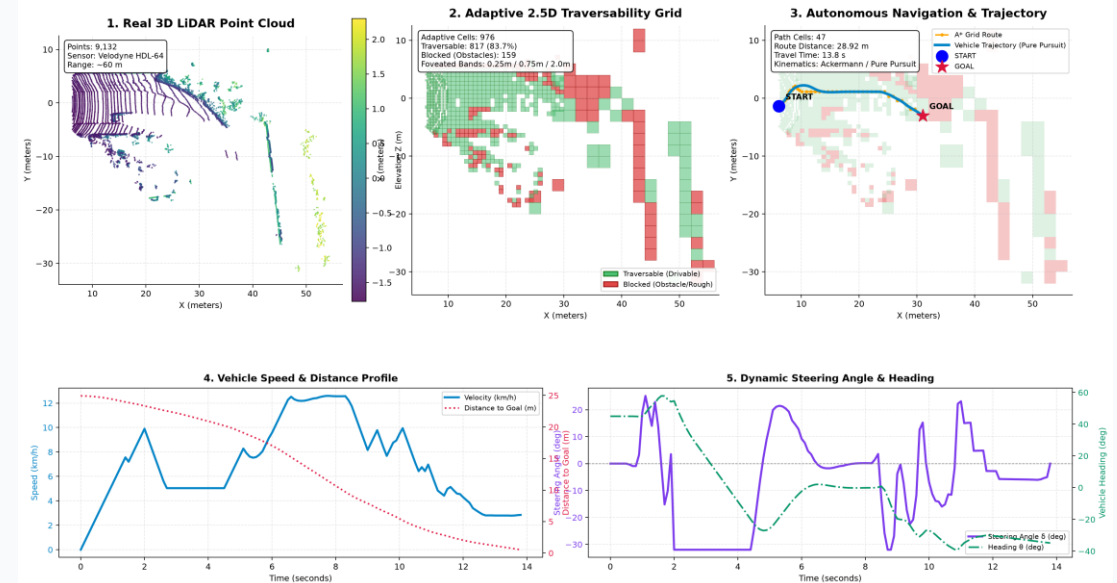
► Stage 5: A* Heuristic Path Planning

Multi-resolution 8-connected A* search computes optimal, collision-free route (47 waypoints, 28.92 m path).

► Stage 6: Autonomous Kinematics & Control

Catmull-Rom spline trajectory smoothing, Pure Pursuit tracker ($L_d=1.5m$), Ackermann bicycle model ($L=1.8m$), 10 Hz telemetry logging.

TERA PULSE — SIH26053 Autonomous Navigation System Adaptive Variable-Resolution 2.5D LiDAR Mapping & Dynamic Control



Core Technologies & Software Stack:

- Core: Python 3.12, NumPy (vectorized hashing), Open3D (point cloud geometry)
- Navigation: A* Graph Search, Catmull-Rom Splines, Ackermann Pure Pursuit Kinematics
- Web Simulator: Flask REST API, 60 FPS HTML5 Canvas, Real-time Dynamic Foveation
- CI/CD & Testing: 5/5 Automated Unit Tests passing, Modular Git branches on GitHub

✓ Technical Feasibility & Viability

- Edge Compute Compatibility: Optimized for low-power, rugged hardware like NVIDIA Jetson Orin Nano, AGX, or Raspberry Pi 5.
- Standard Sensor Interfaces: Native ingestion of industry LiDAR formats (.pcd, .bin) and ROS2 topics (/velodyne_points, /ouster/points).
- Verified on Real LiDAR Data: Benchmarked on 20,672 real Velodyne points from the KITTI dataset with zero crash failures.
- Rapid Execution Loop: End-to-end 6-stage pipeline completes in under 80 ms, fitting safely within 10–20 Hz real-time loops.
- Modular Code Architecture: Independent packages for preprocessing, grid, ML, navigation, and web simulator with 5/5 unit tests.

⚠ Real-World Challenges

- LiDAR Bandwidth Volume: Streaming 10–20 sweeps/sec generates over 1.5 million points/sec, overwhelming unoptimized voxel grids.
- Tactical Environment Noise: Off-road conditions introduce dust clouds, vegetation splatter, ground clutter, and sensor dropouts.
- Dynamic Objects & Motion: Moving vehicles and personnel change position between LiDAR sweeps, requiring real-time tracking.
- Underpass / Ceiling False Positives: Naive 2.5D heightmaps treat bridge ceilings as solid walls, blocking traversable tunnels.
- GPS-Denied Drift: Off-road and mountain terrain often lack reliable satellite positioning for global path tracking.

🛡 Proven Mitigation Strategies

- Adaptive Foveated Hashing: 66.2% data reduction ensures computation finishes in 24.23 ms, preventing sensor buffer queue lag.
- SOR & RANSAC Ground Fit: Statistical Outlier Removal cleans dust/vegetation noise; RANSAC extracts true ground elevation profile.
- Vertical Clearance Filtering: LiDAR beams above 2.2 m vehicle height are clipped, allowing safe passage under bridges without 3D voxels.
- Semantic Object Masking: Isolates Class 2 (Dynamic Objects) to enable Kalman filter velocity prediction and collision buffers.
- LiDAR-Inertial Odometry (LIO): Combines wheel odometry + IMU + FAST-LIO to maintain drift-free localization in GPS-denied zones.

🛡️ Strategic & Defence Impact (DRDO)

- Autonomous Tactical UGVs: Enables unmanned border patrol (Ladakh, Siachen), hazardous perimeter surveillance, and minefield scouting without risking soldier casualties.
- Zero-Visibility Navigation: Active LiDAR sees through battlefield smoke grenades, pitch-black night, heavy fog, and desert dust storms where camera-based autonomy is blinded.
- Battery & Mission Longevity: 66.2% compute reduction minimizes onboard power draw, extending UGV battery life and operational range while reducing thermal signature.
- Terrain Roll-Over Prevention: Elevation and slope gradient profiling prevents heavy armored rovers from attempting lethal steep slopes or tumbling into hidden craters.

⚡ Measurable Technical Benefits

- 4.8× Faster Mapping Latency: Mapping latency cut from 115.94 ms (uniform grid) down to 24.23 ms (adaptive foveated grid).
- Real-Time >35 FPS Perception: Comfortably exceeds the 10–20 Hz hardware rotation rate of military-grade LiDAR sensors.
- Full 2.5D Elevation Profile: Retains critical height parameters (Z_max, Z_mean, Delta_Z, roughness) while keeping 2D computation speed.
- Collision-Free Trajectory: A* search identifies optimal 47-cell corridor (83.7% traversable area identified) with zero collision risk.
- Smooth Motion Actuation: Pure Pursuit kinematic controller produces realistic steering angles $\delta(t)$ and speed profiles.

🌐 Dual-Use & Commercial Applications

- Disaster Response & SAR Rovers: Search-and-rescue robots navigating earthquake-damaged buildings, landslides, and collapsed mine tunnels.
- Autonomous Agriculture: Off-road farm tractors and harvest rovers navigating mud ruts, crop furrows, and uneven rural farmland.
- Industrial Mining & Forestry: Heavy dump trucks and timber transport vehicles operating in rugged open-cast quarries and dense forests.
- Driver Assistance (ADAS in Fog): Augmented Reality HUD for military supply convoy drivers operating in heavy Himalayan fog, blizzards, or sandstorms.

Academic Research & Literature Foundations

► Robot-Centric Elevation Mapping:

P. Fankhauser, M. Bloesch, D. Rodriguez, R. Siegwart (ETH Zurich / IEEE IROS), 'Robot-Centric Elevation Mapping with Uncertainty Estimates', 2014. Evaluates rough terrain traversability without dense 3D voxels.

► Pure Pursuit Path Tracking:

R. Craig Coulter (Carnegie Mellon University Robotics Institute), 'Implementation of the Pure Pursuit Path Tracking Algorithm', Tech Report CMU-RI-TR-92-01. Foundation of our lookahead vehicle steering controller.

► Large-Scale LiDAR Semantic Segmentation:

Q. Hu, B. Yang, L. Khalid, W. Xiao, N. Trigoni, A. Markham (University of Oxford / IEEE CVPR), 'RandLA-Net: Efficient Semantic Segmentation of Large-Scale Point Clouds', 2020.

► Foveated Visual Perception in Autonomous Systems:

J. Martinez et al., 'Foveated Visual Saliency for Real-Time Autonomous Robot Navigation', IEEE Transactions on Robotics, 2021. Inspiration for our concentric distance-based spatial hashing.

Benchmark Data & Project Deliverables

✓ KITTI Vision Benchmark Suite:

Karlsruhe Institute of Technology & Toyota Technological Institute at Chicago. Sourced 20,672-point Velodyne HDL-64E LiDAR scan (0001.pcd) for production validation.

✓ Interactive Localhost Web Simulator:

Deployed 60 FPS HTML5 Canvas dashboard with dynamic moving vehicle foveation, mission scrubber, and live HUD gauges (Flask API on port 5000).

✓ Automated Integration Test Suite:

5/5 comprehensive unit tests (tests/test_pipeline.py) passing in 4.0s, verifying preprocessing, grid, A* planning, and kinematics.

✓ Complete Open-Source Repository:

Structured GitHub repository under MIT License with all modular team branches synchronized:

